

02239 Data Security Privacy

Sebastian Mödersheim



October 9, 2024

Outline

- ① Why Privacy?
- ② Defining Privacy
- ③ Mixes
- ④ Zero-Knowledge Protocols

Outline

- ① Why Privacy?
- ② Defining Privacy
- ③ Mixes
- ④ Zero-Knowledge Protocols

Two Extreme Sides...

- “An honest person has nothing to hide!”



- “1984”



An Example [from Mark Ryan]

- Parents and their teenage daughter:
 - ★ The daughter wants to go out, but not tell the parents where she goes.
 - ★ The parents want to know where she is in case of any emergency.
 - ★ Both are legitimate interests!
 - ★ **Danger:** the parents may overreach – out of concern for their daughter – not respecting her privacy.

An Example [from Mark Ryan]

- Parents and their teenage daughter:
 - ★ The daughter wants to go out, but not tell the parents where she goes.
 - ★ The parents want to know where she is in case of any emergency.
 - ★ Both are legitimate interests!
 - ★ **Danger:** the parents may overreach – out of concern for their daughter – not respecting her privacy.
- Non-electronic solution:
 - ★ The daughter writes where she is going in a sealed letter.
 - ★ The parents can open the letter, but are compelled not to — unless there is an emergency
- Technical solution: with trusted hardware (e.g. a TPM)

Why Privacy?



Why not vote in public?

Why Privacy?

- Being observed, or believing to be observed, can have an influence on ones behavior.
 - ★ Feeling compelled/coerced to act according to others' expectations.
 - ★ Can mean a subtle restriction of the freedom.
- This is getting more sensitive than 20 years ago. Life leaves more traces in different IT systems. Technology allows for:
 - ★ cheap storage of data
 - ★ cheap evaluations of data
 - ★ cheap surveillance
- Protesting against a totalitarian regime, exchange information with other opposition members.

Outline

- ① Why Privacy?
- ② Defining Privacy**
- ③ Mixes
- ④ Zero-Knowledge Protocols

Defining Privacy

What *is* privacy really?

Several informal and semi-formal notions:

- Anonymity: the identity of the actors are protected
- Unlinkability: different actions of the same actor cannot be associated
- ...

Formally Defining Privacy

- Difficult: it is not a classical secrecy problem!
 - ★ Example: poll where you can vote *yes* or *no*
 - ★ *yes* and *no* are values that the intruder knows, also the names of the voters may be no secrets.
 - ★ What is actually secret is: who voted *yes* and who voted *no*!

Formally Defining Privacy

- Difficult: it is not a classical secrecy problem!
 - ★ Example: poll where you can vote *yes* or *no*
 - ★ *yes* and *no* are values that the intruder knows, also the names of the voters may be no secrets.
 - ★ What is actually secret is: who voted *yes* and who voted *no*!
- Inspiration from Cryptography: security formulated as *equivalence* notions:

$$\text{crypt}(k, 0) \sim \text{crypt}(k, 1)?$$

The intruder knows that the plain-text is either 0 or 1, the challenge is for him to tell correctly whether it is 0 or 1.

Formally Defining Privacy

- Difficult: it is not a classical secrecy problem!
 - ★ Example: poll where you can vote *yes* or *no*
 - ★ *yes* and *no* are values that the intruder knows, also the names of the voters may be no secrets.
 - ★ What is actually secret is: who voted *yes* and who voted *no*!
- Inspiration from Cryptography: security formulated as *equivalence* notions:

$$\text{crypt}(k, 0) \sim \text{crypt}(k, 1)?$$

The intruder knows that the plain-text is either 0 or 1, the challenge is for him to tell correctly whether it is 0 or 1.

- Similar in formal verification (with perfect cryptography): static equivalence of frames.

Alpha-Beta Privacy

- Novel approach to formulating privacy:
 - ★ Specify by a logical formula α what information the intruder is allowed to know, e.g. election result:

$$\alpha \equiv v_1, \dots, v_n \in \{0, 1\} \wedge \sum_{i=0}^n v_i = 42$$

- ★ Specify by a logical formula β what the intruder actually knows, e.g., observed cryptographic messages, what he knows about their structure, etc.
 - ★ Privacy: the intruder cannot derive anything from β except what already follows from α .
- Ongoing research effort by Luca Viganò, Sébastien Gondron, Laouen Fernet and myself.
 - ★ We welcome more colleagues to join us e.g. for a MSc thesis!

Example: Basic Authentication Control (BAC)

RFID passport protocol BAC:

$Tag \rightarrow Reader : n_T$

$Reader \rightarrow Tag : \{n_R, n_T, k_R\}_{sk_e(Tag)},$
 $mac(sk_m(Tag), \{n_R, n_T, k_R\}_{sk_e(Tag)})$

$Tag \rightarrow Reader : msg, mac(sk_m(Tag), msg)$

French implementation of the Tag:

$P_{Tag}(Tag) := \nu n_T. send(n_T). receive(x).$
let $(x_1, x_2) = x$
if $mac(sk_m(Tag), x_1) = x_2$
then let $(y_1, y_2, y_3) = decrypt(sk_e(Tag), x_1)$
if $y_2 = n_T$
then send(msg, mac(sk_m(Tag), msg))
else send(errorNonce)
else send(errorMac)

Unlinkability goal: you cannot tell if two given sessions were done by the **same** tag or by **different** tags. Thus, you cannot **link** actions.

Formalizing Unlinkability

- Let Tag be the set of all tags (finite).
- Whenever a tag starts a session:
 - ★ Pick a new variable z (that did not occur before) to represent the name of the tag.
 - ★ Release in α the information that $z \in Tag$.
- The intruder is thus **allowed** to know **that it is a Tag** ...
- After a number of sessions, the intruder thus knows

$$\alpha = z_1 \in Tag \wedge z_2 \in Tag \wedge \dots \wedge z_n \in Tag$$

Formalizing Unlinkability

- Let Tag be the set of all tags (finite).
- Whenever a tag starts a session:
 - ★ Pick a new variable z (that did not occur before) to represent the name of the tag.
 - ★ Release in α the information that $z \in Tag$.
- The intruder is thus **allowed** to know **that it is a Tag** ...
- After a number of sessions, the intruder thus knows

$$\alpha = z_1 \in Tag \wedge z_2 \in Tag \wedge \dots \wedge z_n \in Tag$$

- It is thus a privacy violation, if the intruder ...
 - ★ can find out the precise identity of any tag, e.g. $z_3 = tag_{42}$;
 - ★ can narrow down the identity of any tag, e.g.
 $z_3 \in \{tag_{42}, tag_{112}, tag_{1001}\}$;
 - ★ can say that two tags **are the same**, e.g., $z_3 = z_5$;
 - ★ can say that two tags **are different**, e.g., $z_3 \neq z_5$;

An Attack

$z_1 \rightarrow \text{Reader} : n_T$

$\text{Reader} \rightarrow z_1 : \{\{n_R, n_T, k_R\}_{sk_e(z_1)}, \text{mac}(sk_m(z_1), \{\{n_R, n_T, k_R\}_{sk_e(z_1)}\})\}$

$z_1 \rightarrow \text{Reader} : \text{msg}, \text{mac}(sk_m(z_1), \text{msg})$

$z_2 \rightarrow i(\text{Reader}) : n'_T$

$i(\text{Reader}) \rightarrow z_2 : \{\{n_R, n_T, k_R\}_{sk_e(z_1)}, \text{mac}(sk_m(z_1), \{\{n_R, n_T, k_R\}_{sk_e(z_1)}\})\}$

$z_2 \rightarrow i(\text{Reader}) : \text{errorMac}$

$P_{\text{Tag}}(\text{Tag}) := \nu n_T. \text{send}(n_T). \text{receive}(x).$
 let $(x_1, x_2) = x$
 if $\text{mac}(sk_m(\text{Tag}), x_1) = x_2$
 then let $(y_1, y_2, y_3) = \text{decrypt}(sk_e(\text{Tag}), x_1)$
 if $y_2 = n_T$
 then send(msg, mac(sk_m(Tag), msg))
 else send(errorNonce)
 else send(errorMac)

An Attack

$z_1 \rightarrow \text{Reader} : n_T$

$\text{Reader} \rightarrow z_1 : \{\{n_R, n_T, k_R\}_{sk_e(z_1)}, \text{mac}(sk_m(z_1), \{\{n_R, n_T, k_R\}_{sk_e(z_1)}\})\}$

$z_1 \rightarrow \text{Reader} : \text{msg}, \text{mac}(sk_m(z_1), \text{msg})$

$z_2 \rightarrow i(\text{Reader}) : n'_T$

$i(\text{Reader}) \rightarrow z_2 : \{\{n_R, n_T, k_R\}_{sk_e(z_1)}, \text{mac}(sk_m(z_1), \{\{n_R, n_T, k_R\}_{sk_e(z_1)}\})\}$

$z_2 \rightarrow i(\text{Reader}) : \text{errorMac}$

```
PTag(Tag) :=  $\nu n_T$ .send( $n_T$ ).receive( $x$ ).  
    let ( $x_1, x_2$ ) =  $x$   
    if  $\text{mac}(sk_m(\text{Tag}), x_1) = x_2$   
    then let ( $y_1, y_2, y_3$ ) =  $\text{decrypt}(sk_e(\text{Tag}), x_1)$   
        if  $y_2 = n_T$   
        then send(msg,  $\text{mac}(sk_m(\text{Tag}), \text{msg})$ )  
        else send(errorNonce)  
    else send(errorMac)
```

- If observe **errorMac**: the **mac check** must have failed: $z_1 \neq z_2$.

An Attack

$z_1 \rightarrow \text{Reader} : n_T$

$\text{Reader} \rightarrow z_1 : \{\{n_R, n_T, k_R\}_{sk_e(z_1)}, \text{mac}(sk_m(z_1), \{\{n_R, n_T, k_R\}_{sk_e(z_1)}\})\}$

$z_1 \rightarrow \text{Reader} : \text{msg}, \text{mac}(sk_m(z_1), \text{msg})$

$z_2 \rightarrow i(\text{Reader}) : n'_T$

$i(\text{Reader}) \rightarrow z_2 : \{\{n_R, n_T, k_R\}_{sk_e(z_1)}, \text{mac}(sk_m(z_1), \{\{n_R, n_T, k_R\}_{sk_e(z_1)}\})\}$

$z_2 \rightarrow i(\text{Reader}) : \text{errorMac}$

$P_{\text{Tag}}(\text{Tag}) := \nu n_T. \text{send}(n_T). \text{receive}(x).$
 let $(x_1, x_2) = x$
 if $\text{mac}(sk_m(\text{Tag}), x_1) = x_2$
 then let $(y_1, y_2, y_3) = \text{decrypt}(sk_e(\text{Tag}), x_1)$
 if $y_2 = n_T$
 then send(msg, $\text{mac}(sk_m(\text{Tag}), \text{msg})$)
 else send(errorNonce)
 else send(errorMac)

- If observe **errorMac**: the **mac check** must have failed: $z_1 \neq z_2$.
- If observe **errorNonce**: the **mac check** had worked: $z_1 = z_2$.

The British Implementation

```
 $P_{\text{Tag}}(\text{Tag}) := \nu n_T. \text{send}(n_T). \text{receive}(x).$   
  let  $(x_1, x_2) = x$   
  if  $\text{mac}(sk_m(\text{Tag}), x_1) = x_2$   
  then let  $(y_1, y_2, y_3) = \text{decrypt}(sk_e(\text{Tag}), x_1)$   
    if  $y_2 = n_T$   
    then send(msg, mac( $sk_m(\text{Tag})$ , msg))  
    else send(error)  
  else send(error)
```

The British Implementation


$P_{\text{Tag}}(\text{Tag}) \quad := \quad \nu n_T. \text{send}(n_T). \text{receive}(x).$
 let $(x_1, x_2) = x$
 if $\text{mac}(sk_m(\text{Tag}), x_1) = x_2$
 then let $(y_1, y_2, y_3) = \text{decrypt}(sk_e(\text{Tag}), x_1)$
 if $y_2 = n_T$
 then send(msg, mac($sk_m(\text{Tag})$, msg))
 else send(error)
 else send(error)

- Unlinkable! (as can be shown)


Outline

- ① Why Privacy?
- ② Defining Privacy
- ③ Mixes**
- ④ Zero-Knowledge Protocols


Mixes/Onion Routing




Mixes/Onion Routing




Mixes/Onion Routing




Mixes/Onion Routing




Mixes/Onion Routing




Mixes/Onion Routing




Mixes/Onion Routing



Mixes/Onion Routing




Mixes/Onion Routing



- Each node knows only who is the previous/next on the chain
- Only the **destination** learns the message M
- Nobody knows who is **source** except the source itself
- Nobody knows who is the **destination** (except source and destination)
- Attacker model: can see all exchanged messages and may control **some nodes** (may include **destination**)

Mixes/Onion Routing: Assumptions



- Sender knows the true public key of all nodes.
- At least one node on the path is not controlled by the attacker
- Traffic is evenly distributed between all nodes
- Replay of messages is prevented
- Length of an encrypted message does not reveal the number of encryption layers.
- Note also: if **destination** is not part of mix network, cleartext messages are transmitted on the last leg.

Outline

- ① Why Privacy?
- ② Defining Privacy
- ③ Mixes
- ④ Zero-Knowledge Protocols**

Zero-knowledge proofs

In zero-knowledge proofs we can usually specify a **statement** that is being proved.

- Definitely, that statement is revealed to the verifier
- The verifier (or others) should not learn anything else
- Everybody can draw conclusions from everything they learned

Zero-Knowledge Protocols

Scenario

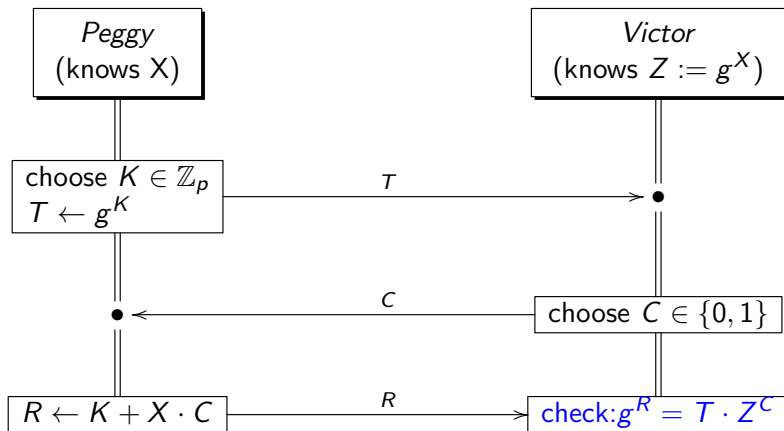
- A **Prover** Peggy and a **Verifier** Victor
- Peggy has a secret and wants to convince Victor of that fact **without** telling the secret.
- Example **Schnorr protocol**:
 - ★ Victor knows a public value $Z \in \mathbb{Z}_p^*$.
 - ★ Peggy wants to prove that she **knows** $X \in \mathbb{Z}_p^*$ with $Z \equiv_p g^X$.
- Why? What's the point of proving this?

Zero-Knowledge Protocols

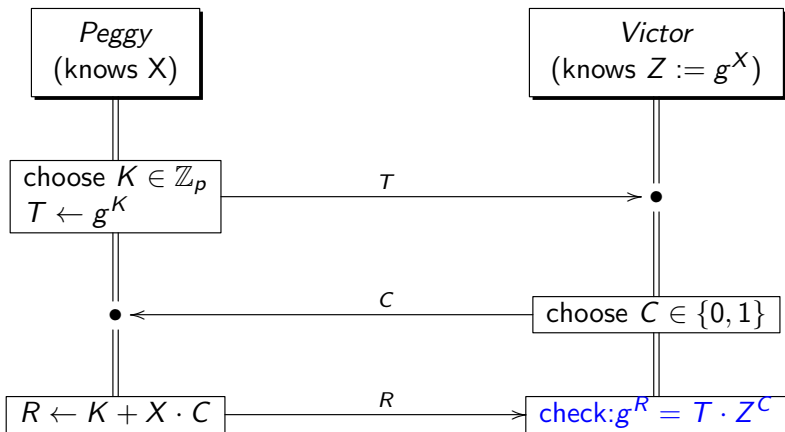
Scenario

- A **Prover** Peggy and a **Verifier** Victor
- Peggy has a secret and wants to convince Victor of that fact **without** telling the secret.
- Example **Schnorr protocol**:
 - ★ Victor knows a public value $Z \in \mathbb{Z}_p^*$.
 - ★ Peggy wants to prove that she **knows** $X \in \mathbb{Z}_p^*$ with $Z \equiv_p g^X$.
- Why? What's the point of proving this?
 - ★ Peggy can authenticate this way
 - ★ Anonymous credential systems ("Private authentication")
 - ★ Voting protocols like Helios/Belenios
 - ★ Alternatives to Proof-of-Work in blockchain implementations.

Schnorr: the Protocol



Schnorr: the Protocol



If Peggy worked correctly:

$$g^R = g^{K+X \cdot C} = g^K \cdot g^{X \cdot C} = T \cdot Z^C$$

Schnorr: Trying to Cheat

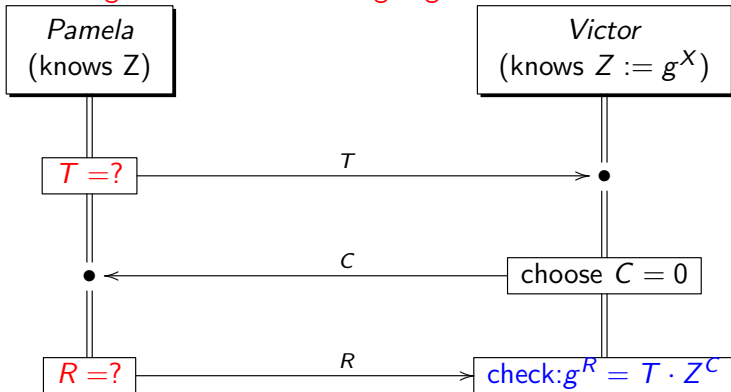
Pamela does not know the secret, but tries to prove its knowledge.

Pamela guesses that Victor is going to choose $C = 0$

Schnorr: Trying to Cheat

Pamela does not know the secret, but tries to prove its knowledge.

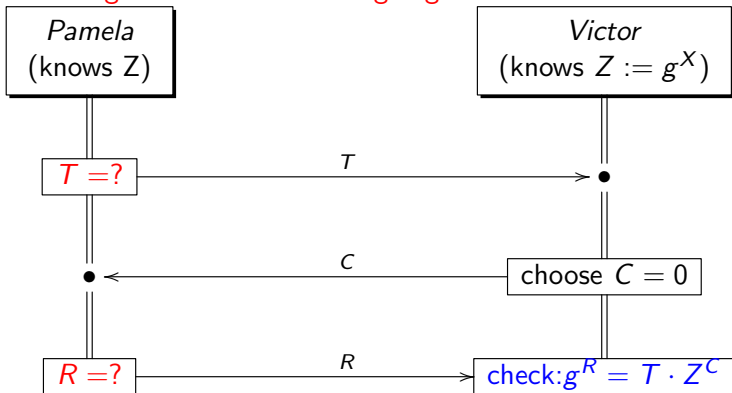
Pamela guesses that Victor is going to choose $C = 0$



Schnorr: Trying to Cheat

Pamela does not know the secret, but tries to prove its knowledge.

Pamela guesses that Victor is going to choose $C = 0$




Strategy for Pamela: Choose R randomly and set $T = g^R$.

Schnorr: Trying to Cheat

Pamela guesses that Victor is going to choose $C = 1$

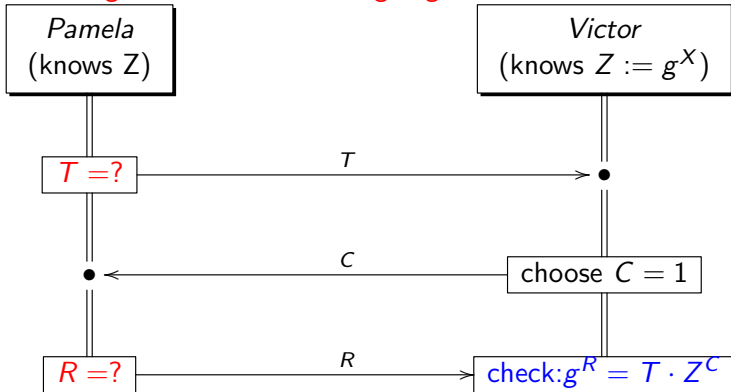
Schnorr: Trying to Cheat

Pamela guesses that Victor is going to choose $C = 1$



Schnorr: Trying to Cheat


Pamela guesses that Victor is going to choose $C = 1$



Find numbers T and R such that $g^R = T \cdot Z$.

Schnorr: Trying to Cheat

Pamela guesses that Victor is going to choose $C = 1$

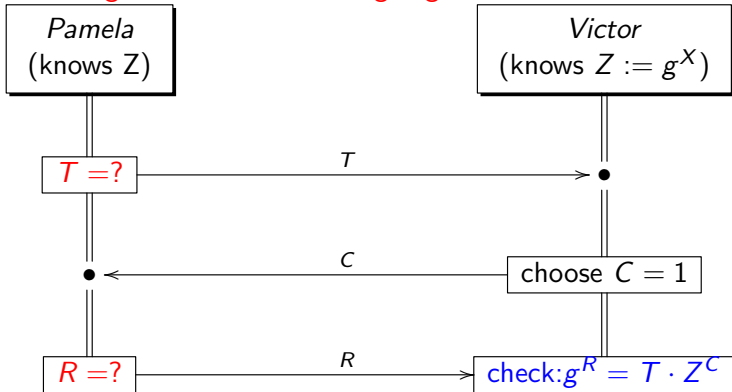


Find numbers T and R such that $g^R = T \cdot Z$.

Hint: division modulo p is also easy!

Schnorr: Trying to Cheat

Pamela guesses that Victor is going to choose $C = 1$



Find numbers T and R such that $g^R = T \cdot Z$.

Hint: division modulo p is also easy!

Strategy for Pamela: Choose R randomly, set $T = \frac{g^R}{Z}$.

Schnorr: Trying to Cheat

- Pamela has a strategy to cheat if $C = 0$:

Choose R randomly, set $T = g^R$.

Schnorr: Trying to Cheat

- Pamela has a strategy to cheat if $C = 0$:

Choose R randomly, set $T = g^R$.

- Pamela has a strategy to cheat if $C = 1$:

Choose R randomly, set $T = \frac{g^R}{Z}$

Schnorr: Trying to Cheat

- Pamela has a strategy to cheat if $C = 0$:

Choose R randomly, set $T = g^R$.

- Pamela has a strategy to cheat if $C = 1$:

Choose R randomly, set $T = \frac{g^R}{Z}$

- Since T must be sent before C is known, she can **prepare** only for one possibility
 - ★ Chance of $\frac{1}{2}$ to cheat (if $C \in \{0, 1\}$ uniformly chosen)

Schnorr: Trying to Cheat

- Pamela has a strategy to cheat if $C = 0$:

Choose R randomly, set $T = g^R$.

- Pamela has a strategy to cheat if $C = 1$:

Choose R randomly, set $T = \frac{g^R}{Z}$

- Since T must be sent before C is known, she can **prepare** only for one possibility
 - ★ Chance of $\frac{1}{2}$ to cheat (if $C \in \{0, 1\}$ uniformly chosen)
- Claim: Pamela cannot do better (to prove...)

Schnorr: Trying to Cheat

- Pamela has a strategy to cheat if $C = 0$:

Choose R randomly, set $T = g^R$.

- Pamela has a strategy to cheat if $C = 1$:

Choose R randomly, set $T = \frac{g^R}{Z}$

- Since T must be sent before C is known, she can **prepare** only for one possibility
 - ★ Chance of $\frac{1}{2}$ to cheat (if $C \in \{0, 1\}$ uniformly chosen)
- Claim: Pamela cannot do better (to prove...)
- If Victor accepts only after n successful rounds of the protocol, the chance to cheat is only 2^{-n} .

Schnorr: Trying to Cheat

Claim

Pamela has no strategy to cheat for unpredictable C .

Proof sketch

Schnorr: Trying to Cheat

Claim

Pamela has no strategy to cheat for unpredictable C .

Proof sketch

- Suppose Pamela has a strategy that works both for $C = 0$ and for $C = 1$.

Schnorr: Trying to Cheat

Claim

Pamela has no strategy to cheat for unpredictable C .

Proof sketch

- Suppose Pamela has a strategy that works both for $C = 0$ and for $C = 1$.
- Then she has values T , R_0 , and R_1 such that
 - ★ (If Victor chooses $C = 0$:) $g^{R_0} = T$
 - ★ (If Victor chooses $C = 1$:) $g^{R_1} = T \cdot Z$

Schnorr: Trying to Cheat

Claim

Pamela has no strategy to cheat for unpredictable C .

Proof sketch

- Suppose Pamela has a strategy that works both for $C = 0$ and for $C = 1$.
- Then she has values T , R_0 , and R_1 such that
 - ★ (If Victor chooses $C = 0$:) $g^{R_0} = T$
 - ★ (If Victor chooses $C = 1$:) $g^{R_1} = T \cdot Z$
- Then she can compute $Q = R_1 - R_0$.

Schnorr: Trying to Cheat

Claim

Pamela has no strategy to cheat for unpredictable C .

Proof sketch

- Suppose Pamela has a strategy that works both for $C = 0$ and for $C = 1$.
- Then she has values T , R_0 , and R_1 such that
 - ★ (If Victor chooses $C = 0$:) $g^{R_0} = T$
 - ★ (If Victor chooses $C = 1$:) $g^{R_1} = T \cdot Z$
- Then she can compute $Q = R_1 - R_0$.
- It holds that $g^Q = g^{R_1 - R_0} = \frac{g^{R_1}}{g^{R_0}} = \frac{T \cdot Z}{T} = Z$.

Schnorr: Trying to Cheat

Claim

Pamela has no strategy to cheat for unpredictable C .

Proof sketch

- Suppose Pamela has a strategy that works both for $C = 0$ and for $C = 1$.
- Then she has values T , R_0 , and R_1 such that
 - ★ (If Victor chooses $C = 0$:) $g^{R_0} = T$
 - ★ (If Victor chooses $C = 1$:) $g^{R_1} = T \cdot Z$
- Then she can compute $Q = R_1 - R_0$.
- It holds that $g^Q = g^{R_1 - R_0} = \frac{g^{R_1}}{g^{R_0}} = \frac{T \cdot Z}{T} = Z$.
- So with $Q = X$ she indeed knows the discrete logarithm of Z .

Schnorr: Trying to Cheat

Claim

Pamela has no strategy to cheat for unpredictable C .

Proof sketch

- Suppose Pamela has a strategy that works both for $C = 0$ and for $C = 1$.
- Then she has values T , R_0 , and R_1 such that
 - ★ (If Victor chooses $C = 0$:) $g^{R_0} = T$
 - ★ (If Victor chooses $C = 1$:) $g^{R_1} = T \cdot Z$
- Then she can compute $Q = R_1 - R_0$.
- It holds that $g^Q = g^{R_1 - R_0} = \frac{g^{R_1}}{g^{R_0}} = \frac{T \cdot Z}{T} = Z$.
- So with $Q = X$ she indeed knows the discrete logarithm of Z .
- **So it is not cheating!**

Schnorr: Curious Victor

Victor would like to learn the secret X ...

Zero-Knowledge Property

Victor learns nothing except the proved statement.

Proof sketch

Schnorr: Curious Victor

Victor would like to learn the secret X ...

Zero-Knowledge Property

Victor learns nothing except the proved statement.

Proof sketch

Consider a transcript of an exchange between Peggy and Victor. (Assume Victor chooses challenges randomly.)

$$\begin{array}{ccc} T_1 & C_1 & R_1 \\ T_2 & C_2 & R_2 \\ & \dots & \\ T_n & C_n & R_n \end{array}$$

Schnorr: Curious Victor

Victor would like to learn the secret X ...

Zero-Knowledge Property

Victor learns nothing except the proved statement.

Proof sketch

Consider a transcript of an exchange between Peggy and Victor. (Assume Victor chooses challenges randomly.)

Create a fake-transcript using random challenges C'_i

T_1 C_1 R_1

T_2 C_2 R_2

...

T_n C_n R_n

C'_1

C'_2

...

C'_n

Schnorr: Curious Victor

Victor would like to learn the secret X ...

Zero-Knowledge Property

Victor learns nothing except the proved statement.

Proof sketch

Consider a transcript of an exchange between Peggy and Victor. (Assume Victor chooses challenges randomly.)

$$\begin{array}{ccc} T_1 & C_1 & R_1 \\ T_2 & C_2 & R_2 \\ & \dots & \\ T_n & C_n & R_n \end{array}$$

Create a fake-transcript using random challenges C'_i and then filling the T'_i and R'_i according to the cheat strategies of Pamela for known challenges:

$$\begin{array}{ccc} T'_1 & C'_1 & R'_1 \\ T'_2 & C'_2 & R'_2 \\ & \dots & \\ T'_n & C'_n & R'_n \end{array}$$

Schnorr: Curious Victor

Victor would like to learn the secret X ...

Zero-Knowledge Property

Victor learns nothing except the proved statement.

Proof sketch

- Give the two transcripts to a third party. Can one tell the real from the fake?

Schnorr: Curious Victor

Victor would like to learn the secret X ...

Zero-Knowledge Property

Victor learns nothing except the proved statement.

Proof sketch

- Give the two transcripts to a third party. Can one tell the real from the fake?
- **No**: statically indistinguishable. (Not even distinguishable with unbounded computing resources!)

Schnorr: Curious Victor

Victor would like to learn the secret X ...

Zero-Knowledge Property

Victor learns nothing except the proved statement.

Proof sketch

- Give the two transcripts to a third party. Can one tell the real from the fake?
- **No**: statically indistinguishable. (Not even distinguishable with unbounded computing resources!)
- So, whatever information Victor is getting out of the transcript, he may have created that himself without Peggy!

The Mafia-owned Restaurant

Claim (Shamir)

I can go to a Mafia-owned restaurant and pay with my zero-knowledge proof-based credit card frequently, and yet the mobsters cannot impersonate me.


The Mafia-owned Restaurant

Claim (Shamir)

I can go to a Mafia-owned restaurant and pay with my zero-knowledge proof-based credit card frequently, and yet the mobsters cannot impersonate me.

They can, if they do it online.

The Mafia Attack



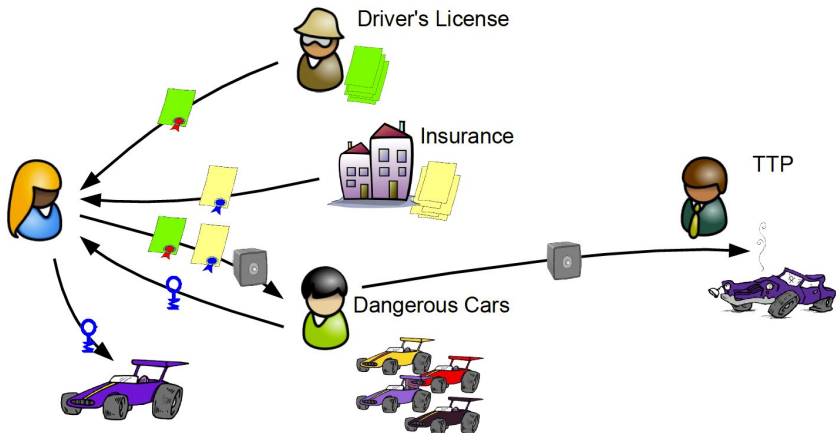
Classical Man in the middle attack: forward all messages!

Countermeasures

- Distance bounding protocols.
- Link the intended verifier to the statement being proved.

IBM's Idemix

A privacy friendly credential system



IBM's Idemix

- Credentials: list of attribute-value pairs, signed by an issuer.
`sign(copenhagen-commune; name="Johanna Gossner",
birthdate=14/12/73, kind="Drivers License", class="3", ...)`
- Signatures are special: Camenisch-Lysyanskaya. Allows for special zero-knowledge proofs
 - ★ Proving possession of credential without transmitting it
 - ★ Revealing only some attributes (e.g. kind and class)
 - ★ Proving properties about attributes (e.g. $\text{birthdate} < 9/10/2000$)
 - ★ Proving relations between credentials (e.g. passport on same name)
 - ★ Signing a statement (e.g. "I confirm this electronic order...")

IBM's Idemix

- Credentials: list of attribute-value pairs, signed by an issuer.
`sign(copenhagen-commune; name=" Johanna Gossner",
birthdate=14/12/73, kind=" Drivers License", class=" 3", ...)`
- Signatures are special: Camenisch-Lysyanskaya. Allows for special zero-knowledge proofs
 - ★ Proving possession of credential without transmitting it
 - ★ Revealing only some attributes (e.g. kind and class)
 - ★ Proving properties about attributes (e.g. $\text{birthdate} < 9/10/2000$)
 - ★ Proving relations between credentials (e.g. passport on same name)
 - ★ Signing a statement (e.g. "I confirm this electronic order...")
- Privacy: A verifier cannot find out more than what was proved.
 - ★ E.g. different transactions are unlikable (unless linked by prover).
 - ★ Holds even if the issuer is dishonest and collaborates with a dishonest verifier.

Idemix: Privacy with Accountability?

How can we ensure that we

- protect the privacy of the innocent
- but not the privacy of the guilty?

Idemix: Privacy with Accountability?

How can we ensure that we

- protect the privacy of the innocent
 - but not the privacy of the guilty?
-
- Can be achieved using escrow with [verifiable encryptions](#)

Idemix: Privacy with Accountability?

How can we ensure that we

- protect the privacy of the innocent
 - but not the privacy of the guilty?
-
- Can be achieved using escrow with **verifiable encryptions**
 - Provers have to encrypt their real names with the public key of an authority.

Idemix: Privacy with Accountability?

How can we ensure that we

- protect the privacy of the innocent
 - but not the privacy of the guilty?
-
- Can be achieved using escrow with **verifiable encryptions**
 - Provers have to encrypt their real names with the public key of an authority.
 - Provers have to prove (in zero-knowledge) that they have encrypted the real name (and it is correctly encrypted).

Idemix: Privacy with Accountability?

How can we ensure that we

- protect the privacy of the innocent
 - but not the privacy of the guilty?
-
- Can be achieved using escrow with **verifiable encryptions**
 - Provers have to encrypt their real names with the public key of an authority.
 - Provers have to prove (in zero-knowledge) that they have encrypted the real name (and it is correctly encrypted).
 - Then the authority can **revoke** their privacy when needed.

Idemix: Privacy with Accountability?

How can we ensure that we

- protect the privacy of the innocent
 - but not the privacy of the guilty?
-
- Can be achieved using escrow with **verifiable encryptions**
 - Provers have to encrypt their real names with the public key of an authority.
 - Provers have to prove (in zero-knowledge) that they have encrypted the real name (and it is correctly encrypted).
 - Then the authority can **revoke** their privacy when needed.
 - This should be tied to a legal system, e.g. the privacy revocation happens only on court order.